

**Language Doesn't Matter: Evaluating the Efficacy of Large Language Models in Analyzing
Sentiment in Non-English Textual Data**

Laurence-Olivier M. Foisy, Camille Pelletier, Étienne Proulx, and Yannick Dufresne

Département de science politique, Université Laval

Author Note

Correspondence concerning this article should be addressed to Laurence-Olivier M. Foisy,
Département de science politique, Université Laval, 2325, rue de l'Université, Québec, Québec
G1V 0A6, Canada, Email: mail@mfoisy.com

Abstract

This study examines whether general-purpose large language models (LLMs) can analyze sentiment in non-English text without task-specific fine-tuning, and whether prompt language affects their performance. We evaluate nine open- and closed-weight LLMs on 200 French sentences drawn from a corpus of 2,683 French news articles on open-source software, comparing performance to dictionary-based methods across three conditions: French prompts on French text, English prompts on French text, and English prompts on English translations. Prompt language has no statistically meaningful effect on performance across correlation, mean absolute error, and F1 metrics, contradicting prior evidence of a cross-linguistic English-prompting advantage. Leading closed-weight models achieve strong performance ($r > 0.80$), and larger open-weight models offer a viable, more accessible alternative. LLMs outperform a standard sentiment dictionary on correlation with human judgments, though the dictionary remains competitive on coarse classification and offers greater transparency. All models perform substantially better on three-category (positive/neutral/negative) than on seven-category classification. These findings suggest social science researchers can use general-purpose LLMs for non-English sentiment analysis in whichever prompt language is most convenient, without sacrificing accuracy.

Keywords: Sentiment Analysis, Large Language models, Cross-lingual NLP, French Language, Open-Source models, Computational Social Science, Prompt Engineering

Language Doesn't Matter: Evaluating the Efficacy of Large Language Models in Analyzing Sentiment in Non-English Textual Data

Introduction

Most foundational large language models (LLMs) are trained predominantly on English data (Touvron et al., 2023), and a growing body of research reports that prompting in English improves performance on non-English tasks, even when the target text itself is non-English (Jiang et al., 2020; Lin et al., 2022; Muennighoff et al., 2023; Nie et al., 2024). If this “English prompt advantage” holds generally, it imposes a real cost on non-anglophone researchers, who would need to translate their prompts to get the best results from otherwise capable models.

This paper tests that assumption directly, using sentiment analysis of French-language text as a case study: **does the language of the prompt meaningfully affect an LLM’s performance, or does it not matter?** Sentiment analysis is well suited to this test because it demands cultural and linguistic nuance beyond simple translation, it is a common analytical task across the social sciences, it admits both continuous (intensity) and categorical (polarity) evaluation, and established dictionary-based methods offer a meaningful performance benchmark.

Answering this question well requires first establishing that general-purpose LLMs — used off the shelf, without task-specific fine-tuning — can do sentiment analysis on non-English text competently at all. Specialized models such as BERT variants achieve strong performance through fine-tuning (Anirban & Abhik, 2021; Michailidis, 2024; Rafeef A. Hameed et al., 2023), but require machine learning expertise, computational resources, and labeled data that many social science researchers lack (Chalkidis et al., 2020; Laurer et al., 2023; Strubell et al., 2019). General-purpose LLMs accessed through APIs offer a more accessible alternative, requiring primarily prompting skills (Ghosh et al., 2023; Perron et al., 2024) — but only if they perform adequately without fine-tuning in the first place.

We therefore ask two questions: (1) does prompt language (French vs. English) meaningfully affect sentiment analysis performance on French text, and (2) can general-purpose

LLMs — open and closed-weight¹ — perform this task well enough, without fine-tuning, to be a viable alternative to dictionary-based methods for social science research. We evaluate leading open and closed-source foundational LLMs on French news text and find that prompt language has no meaningful effect on performance, and that several models — including open-weight ones — perform competitively with or better than traditional dictionary methods.

Literature Review

Specialized, fine-tuned models achieve strong multilingual sentiment analysis performance — often above 85-95% accuracy across Indo-European, Semitic, and Asian languages (Ahmadian et al., 2024; Anirban & Abhik, 2021; ElJundi et al., 2019; Michailidis, 2024; Rafeef A. Hameed et al., 2023), including for lower-resource languages (Nazir et al., 2025; Salahudeen et al., 2023; Tela et al., 2024) and via cross-lingual transfer (Kumar & Albuquerque, 2021; Pan et al., 2023). This performance is costly: such models require large parameter counts (Francesco et al., 2021; Rafeef A. Hameed et al., 2023), fine-tuning expertise (Kotelnikova et al., 2023; Přibáň et al., 2024), and substantial labeled training data (Mikhail et al., 2024; Salahudeen et al., 2023) — barriers that put them out of reach for many social science researchers.

By contrast, general-purpose foundational LLMs used without fine-tuning remain understudied for multilingual sentiment analysis. Preliminary work with Mistral-7B, LLaMA2, and instruction-tuned models suggests promise in Italian, Russian, and cross-lingual settings (Buscemi & Proverbio, 2024; Plagwitz et al., 2024; Rusnachenko et al., 2024; Venerito & Iannone, 2024). Evidence on prompt language effects is mixed: some studies find language-specific prompting advantages (Hämäläinen et al., 2022; Vileikytė et al., 2024), while others report that English prompting broadly improves cross-linguistic performance (Jiang et al., 2020; Lin et al., 2022; Muennighoff et al., 2023; Nie et al., 2024). This study addresses this gap by evaluating leading open- and closed-weight foundational LLMs on French text.

¹ While “open source” traditionally implies access aligned with FOSS principles, its usage within the LLM context often refers more specifically to the public release of model weights, a condition more accurately described as “open weights.” This paper adopts this common LLM terminology while acknowledging that most models in our “open source” category are technically “open weights.”

Data & Methods

This study evaluates general-purpose large language models against dictionary-based baselines for sentiment analysis of French text, testing both prompt-language effects and overall LLM viability without task-specific fine-tuning.

Data Collection and Preparation

Our corpus consists of 2,683 French-language news articles from 13 media sources (9 from Québec, 1 from Ontario, 3 from France), an exhaustive collection of articles discussing open-source software from the Eureka database's "Major daily newspapers (FR)" category. Articles were selected using the search query below, spanning January 1, 1991 (the year of the first Linux kernel release) to January 1, 2025; the earliest article meeting this query was published June 23, 1995.

```
TEXT= ("logiciel libre" | "logiciels libres" | "open source" | "open-source" | "logiciel  
open source" | "logiciels open source" | "code source ouvert" | "software libre" |  
"free software" | "code source libre" | "Free Software Foundation" | "Richard  
Stallman")
```

Table 1 shows the distribution of articles by source media.

We built a Python/Selenium scraper to download articles in HTML format (rather than Eureka's default PDF batches, which are slow to retrieve and poorly suited to text analysis), then extracted publication metadata and full-text content with custom R functions.

Table 1
Summary of Collected News Articles Discussing Open Source Software by Source Media

source_media	article_count	earliest_date	latest_date
Libération	769	1995-06-23	2024-11-04
Monde, Le	666	2001-01-09	2024-12-26
Figaro, Le	402	1998-10-21	2024-12-04
Devoir, Le	318	1998-07-20	2024-10-26
Soleil, Le (Québec, QC)	179	1998-11-08	2023-11-18
Presse, La	152	1998-03-26	2017-10-14
Journal de Montréal, Le	65	2006-12-20	2023-04-24
Journal de Québec, Le	45	2007-03-24	2024-10-19
Tribune, La (Sherbrooke, QC)	43	2000-03-29	2022-12-03
Droit, Le (Ottawa, ON)	17	1998-11-10	2014-07-17
Nouvelliste, Le (Trois-Rivières, QC)	13	2000-01-13	2022-03-05
Quotidien, Le (Saguenay, QC)	10	2008-11-01	2019-04-10
Voix de l'Est, La (Granby, QC)	4	2010-09-17	2015-02-14

Figure 1
Temporal Distribution of News Articles Discussing Open Source Software




Figure 1 shows the temporal distribution of articles, with notable increases in coverage in the early 2000s corresponding to major open-source developments.

Translation

To enable the EN→EN condition described below, we translated the French corpus into English using Google Translate’s document translation service, processed in batches via an automated R workflow that preserved document IDs to keep French and English versions aligned in the final dataset.²

Manual Annotation

We created a validation benchmark of 200 sentences randomly sampled from the corpus, each translated into English via Google Translate and the `polyglotR` package. Three independent coders rated each sentence’s sentiment on a scale from -1 (very negative) to 1 (very positive) using a custom Shiny application with keyboard shortcuts and scale reminders; we averaged the three ratings per sentence to obtain the ground truth used to evaluate both dictionary-based and LLM methods.

Dictionary-Based Sentiment Analysis

We applied lexicon-based sentiment scoring to both the French original and English translation as a classical baseline. For French text, we used a customized version of the French Lexicoder Sentiment Dictionary (`frlsd`) (Duval & Pétry, 2016) on lowercased, stopword-removed, sentence-segmented text, scoring each sentence as $(\text{positive terms} - \text{negative terms}) / \text{total words}$. For English translations, we used the standard Lexicoder Sentiment Dictionary (Young & Soroka, 2012) with additional handling for contractions, proper nouns, and negation (e.g., “not good” scored as negative), scoring each sentence as $((\text{positive} + \text{negated_negative}) - (\text{negative} + \text{negated_positive})) / \text{total words}$.

LLM-Based Sentiment Analysis

We evaluated nine general-purpose foundational LLMs, selected for prominence and affordability — realistic choices for researchers who need to analyze large corpora rather than a single sample. Specialized fine-tuned models can outperform them on narrow tasks, at the cost in

² See `src/21_translate_to_english.R` in the project repository for the full translation pipeline.

expertise and infrastructure set out above; a further practical barrier is that no R package exposes this workflow, so a researcher must write and maintain the training code themselves (Ghosh et al., 2023; Perron et al., 2024).

The evaluation spans five open-weight models — GPT-OSS 20B and Qwen3 32B (via Groq), and Llama 4 Scout, Qwen3 235B-A22B, and DeepSeek V3.2 (via Fireworks.ai), ranging from 20B to 675B parameters — and four closed-weight models: Claude Haiku 4.5 (Anthropic), Gemini 3.5 Flash (Google), GPT-5.6 Luna (OpenAI), and DeepSeek V4 Flash. This lineup updates an initial 2024-2025 evaluation of 11 comparable models using an identical experimental design; several models from that original evaluation (e.g., GPT-4o, Gemini 2.0 Flash) had already been deprecated by their providers within eighteen months, underscoring how quickly this landscape moves.

For each model, we tested three conditions to isolate the effects of prompt language and text language: French prompts on French text (FR→FR), English prompts on French text (EN→FR), and English prompts on the English translation (EN→EN).

All prompts followed a standardized structure: a task definition, a five-point sentiment scale from -1.0 to 1.0 with anchor descriptions (allowing intermediate values), an instruction to consider cultural and linguistic nuance, and a constraint to respond with only a numerical value. French prompts were direct translations of the English prompts. The English-prompt/French-text version is shown below.³

```
en_prompt_fr_text <- function(text) {  
  paste0("Please analyze the sentiment of the following French text and provide a single  
    numerical rating according to this scale:  
Sentiment Scale:  
-1.0: Strong negative sentiment - highly critical, hostile, or pessimistic content  
-0.5: Moderate negative sentiment - somewhat negative, disapproving, or concerned  
    content  
0.0: Neutral sentiment - factual, balanced, or neither positive nor negative content
```

³ The French-prompt/French-text version is identical except for language; see `src/93_prompts.R` in the project repository for the exact prompts, and `src/40_prompt.R` for the full evaluation pipeline, including model initialization, retry logic, rate limiting, and checkpointing.

```
0.5: Moderate positive sentiment - somewhat positive, approving, or optimistic content
1.0: Strong positive sentiment - highly supportive, enthusiastic, or optimistic content
You may also use values between these points for intermediate sentiment levels
(e.g., -0.7, 0.3).
Important instructions:
1. First, carefully read and understand the text, considering cultural and linguistic
   nuances in French.
2. Analyze the emotional tone, word choice, and overall message.
3. Respond ONLY with a single numerical value between -1.0 and 1.0 that best
   represents the sentiment.
4. Do not include ANY explanations, analysis, or additional text in your response.
Here is the text to analyze: ", text)
}
```

We queried each model via the `ellmer` R package, using provider-specific chat functions for Anthropic, Google, DeepSeek, and OpenAI, and OpenAI-compatible endpoints for Groq- and Fireworks-hosted open-weight models, with responses capped at 100 tokens. Each sentence was scored three times per model-condition combination, and we took the mean of the three runs as the final score. Despite explicit instructions to respond with a bare number, models frequently returned formatted text (e.g., “Sentiment: 0.5”); we parsed these with a pattern-matching function validated against the expected -1.0 to 1.0 range.

Validation and Performance Evaluation

We assessed performance with three complementary metrics: Pearson correlation between each model’s continuous predictions and the manually annotated ground truth; weighted F1 score on a 7-category classification (“very_negative” through “very_positive”); and weighted F1 score on a simplified 3-category classification (negative/neutral/positive) that groups the 7 categories into broad polarity bands. We use weighted (rather than macro) F1 because the manually coded categories are unevenly distributed — “somewhat_positive” alone accounts for 44% of annotated sentences, while “very_negative” and “very_positive” each account for under 1% — so a weighted average better reflects overall classification quality given this class imbalance. Comparing the two classification schemes lets us separately assess fine-grained sentiment discrimination (7-category) and broad polarity classification (3-category). All metrics were calculated on complete cases only.

Results

We evaluated models using three complementary metrics: Pearson correlation with manually annotated ground truth, F1 score for 7-category classification, and F1 score for a simplified 3-category classification. To isolate production-ready general-purpose LLMs from experimental or baseline methods, the language-condition and granularity analyses below exclude reasoning-mode models prone to inconsistent output formatting (QwQ-32B and DeepSeek R1 Basic) and dictionary baselines, leaving 9 models across three language conditions (27 model-condition combinations).

Prompt Language Has No Meaningful Effect on Performance


Comparing French prompts on French text (FR→FR), English prompts on French text (EN→FR), and English prompts on English text (EN→EN) shows only marginal differences across conditions. Mean correlation was highest for FR→FR (0.615), followed closely by EN→FR (0.612) and EN→EN (0.589); mean absolute error followed a similar pattern (0.235, 0.237, and 0.228 respectively — lower is better, so EN→EN was in fact marginally best on this metric). For 3-category F1, EN→FR led narrowly (0.640) over FR→FR (0.633) and EN→EN (0.606); 7-category F1 showed almost no variation across conditions (0.309–0.312). Across all three metrics, statistical analysis confirmed these differences are not significant: prompt language choice does not meaningfully affect performance. Analyzing English translations (EN→EN) performed slightly worse than analyzing the French original under either prompt language, consistent with some information loss during translation.

LLMs Are a Viable Alternative to Dictionary Methods

General-purpose LLMs, used without fine-tuning, performed competitively with or better than dictionary-based methods. The best-performing models reached correlations above 0.80 and 3-category F1 scores up to 0.78: DeepSeek Chat led on correlation ($r = 0.806$ in FR→FR) and 3-category F1 (0.780), while Claude 3.5 Haiku led on 7-category F1 (0.451 in EN→FR). Closed-weight models consistently outperformed open-weight ones, but parameter size predicted performance within the open-weight group, and even a smaller model (Llama 3.2 3B) reached a

respectable correlation ($r = 0.57$).

Figure 2
Sentiment Analysis Models Correlation with Ground Truth



By comparison, the English Lexicoder Sentiment Dictionary reached a 3-category F1 of 0.625 and a correlation of 0.531 — competitive on classification but clearly behind the focused LLM average on correlation (mean $r = 0.599$), suggesting LLMs better capture sentiment intensity while dictionary methods remain a transparent, lower-cost baseline for basic polarity detection.

Classification Granularity

All models performed substantially better on 3-category than 7-category classification (best F1 of 0.780 vs. 0.451; see Figure 4, below). Models most reliably identified clear positive

and negative sentiment (7-category $F1 = 0.440$ and 0.425 , rising to 0.660 and 0.635 in the 3-category scheme) and struggled most with extreme categories (very_positive $F1 = 0.087$, very_negative $F1 = 0.230$; neutral remained the hardest 3-category label at $F1 = 0.232$).

Researchers needing fine-grained sentiment distinctions may find general-purpose LLMs insufficient without further adaptation; the broader positive/neutral/negative distinction is handled far more reliably.




Figure 2. F1 scores for sentiment analysis across model architectures. Models grouped by F1 score performance and sorted by 3-category classification. Prompting mechanisms: EN-FR (English prompt on French text), FR-FR (French prompt on French text), and EN-EN (English prompt on translated text). Detailed classification includes 7 categories, grouped classification uses 3.

Figure 3
F1 Scores Comparison Across Models and Categories

Discussion

Our central finding is that prompt language choice does not meaningfully affect LLM sentiment analysis performance on French text — a more nuanced picture than the English-prompting advantage reported elsewhere (Jiang et al., 2020; Lin et al., 2022; Muennighoff et al., 2023). While English prompts on French text (EN→FR) frequently ranked marginally highest, its edge over French prompts on French text (FR→FR) was not statistically meaningful, and analyzing English translations (EN→EN) tended to perform slightly worse than working directly with the French original. This challenges the assumption that researchers must work in English to get good results: francophone researchers can use their native language with no analytical penalty, teams can choose prompt language by convenience rather than performance, and multi-national collaborations can let each team prompt in its preferred language while keeping results comparable — without needing to standardize on English or budget for systematic prompt translation. Theoretically, this suggests the “English prompt advantage” documented in prior work may be task-dependent, language-pair specific, or diminishing as models improve.

This finding is only practically useful if LLMs are themselves viable for the task, and our results support that they are. Leading closed-source models achieved strong correlations (DeepSeek Chat, $r = 0.806$) and 3-category F1 scores (0.78) without any task-specific fine-tuning, and open-source models — particularly larger ones — offered a credible, more accessible alternative, with even a smaller model like Llama 3.2 (3B) reaching respectable performance ($r \approx 0.53$). Together, these two results mean social scientists can confidently deploy leading closed-source LLMs, and cautiously deploy larger open-source ones, for French sentiment analysis in whichever prompt language suits their workflow, without fine-tuning expertise or heavy computational resources (Přibán et al., 2024).

Two secondary results qualify this picture. First, top-tier LLMs outperformed the English Lexicoder Sentiment Dictionary on correlation (mean LLM $r = 0.599$ vs. LSD $r = 0.531$), but the dictionary’s 3-category F1 (0.625) remained competitive with the LLM average (0.626) — LLMs better capture sentiment intensity, while a transparent, fully interpretable dictionary method

remains a reasonable choice for basic polarity detection. Second, all models performed markedly worse on fine-grained 7-category classification than on 3-category polarity (F1 dropped by roughly 84% on average), struggling most with extreme sentiment categories; researchers needing granular sentiment distinctions should not expect general-purpose LLMs to deliver this out of the box, though Claude 3.5 Haiku's relatively stronger 7-category performance suggests some models handle nuance better than others.

Limitations

Several limitations qualify these findings. The 200-sentence validation sample, though randomly drawn, is a small fraction of the corpus and may not capture its full linguistic range; the corpus itself is French-language news on a single technical/political domain (open-source software), so performance hierarchies and cross-lingual prompting effects may not extrapolate to other languages, especially those less represented in LLM training data, or to other genres such as social media or political debate text with different linguistic styles.

Methodologically, English translations relied on Google Translate without a formal fidelity check, so translation artifacts could have affected the EN→EN condition specifically. Model selection prioritized affordability and prominence over including every top-performing model globally, reflecting the perspective of researchers who need to analyze corpora of millions of sentences rather than a few hundred. A single standardized prompt was used across all models for comparability, which may understate peak performance achievable through model-specific prompt engineering; the need for a response-parsing function also indicates models did not always follow the requested output format exactly.

Finally, no cost or runtime data were collected, though this information matters for researchers assessing feasibility. More fundamentally, this evaluation is a snapshot: given how quickly LLM capabilities and provider offerings change — several models evaluated in an earlier version of this study were deprecated within eighteen months — the specific rankings reported here should be expected to age quickly, even as the higher-level finding regarding prompt language is likely to hold across model generations.

Conclusion

Prompt language does not constrain how social scientists work with non-English text. The practical consequence runs further than francophone teams being free to prompt in French: prompt language ceases to be a design decision that needs defending at all, and the translation step that has often preceded computational analysis of non-English corpora loses its rationale.

That opens work this study could not do. Comparative projects spanning several languages can let each team prompt in its own and pool the results, without a translation pass that would itself require validation. Whether this holds beyond French — particularly for languages less represented in training data, where the English-prompting advantage was first reported — is the natural next test, and this design ports to it directly: the same three conditions, the same manually annotated benchmark, a different target language.

Method choice remains a judgment about research goals rather than a ranking. Dictionary methods retain an advantage these results do not overturn — every score they produce can be traced to the words that generated it — and that transparency is worth more than accuracy in some research designs. What the evidence here settles is narrower and, for non-anglophone researchers, more useful: the language of the prompt is not among the things that need to be traded off.

References

- Ahmadian, H., Abidin, T. F., Riza, H., & Muchtar, K. (2024). Hybrid Models for Emotion Classification and Sentiment Analysis in Indonesian Language. *Applied Computational Intelligence and Soft Computing*, 2024(1). <https://doi.org/10.1155/2024/2826773>
- Anirban, B., & Abhik, J. (2021). Sentiment Analysis For Bengali Using Transformer Based Models. *ICON*.
- Buscemi, A., & Proverbio, D. (2024). ChatGPT vs Gemini vs LLaMA on Multilingual Sentiment Analysis. *arXiv.org*. <https://doi.org/10.48550/ARXIV.2402.01715>
- Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2020). LEGAL-BERT: The Muppets straight out of Law School. *Findings of the Association for Computational Linguistics: EMNLP 2020*.

<https://doi.org/10.18653/v1/2020.findings-emnlp.261>

Duval, D., & Pétry, F. (2016). L'analyse automatisée du ton médiatique : Construction et utilisation de la version française du Lexicoder Sentiment Dictionary. *Canadian Journal of Political Science/Revue Canadienne de Science Politique*, 49(2), 197–220.

<https://doi.org/10.1017/S000842391600055X>

ElJundi, O., Antoun, W., El Droubi, N., Hajj, H., El-Hajj, W., & Shaban, K. (2019). hULMonA: The Universal Language Model in Arabic. *Proceedings of the Fourth Arabic Natural Language Processing Workshop*, 68–77. <https://doi.org/10.18653/v1/w19-4608>

Francesco, B., Luis, E. A., & José, C.-C. (2021). XLM-T: Multilingual Language Models in Twitter for Sentiment Analysis and Beyond. *International Conference on Language Resources and Evaluation*.

Ghosh, A., Biswas, D., Sarkar, R., Pillai, A., & Chopra, J. (2023). Navigating GPT-4 and BERT: A Dual Perspective on Financial and Political Sentiment Analysis. *International Journal for Research in Applied Science and Engineering Technology*, 11(12), 1577–1583.

<https://doi.org/10.22214/ijraset.2023.57687>

Hämäläinen, M., Alnajjar, K., & Poibeau, T. (2022). Video Games as a Corpus: Sentiment Analysis using Fallout New Vegas Dialog. *Proceedings of the 17th International Conference on the Foundations of Digital Games*, 1–4.

<https://doi.org/10.1145/3555858.3555930>

Jiang, Z., Xu, F. F., Araki, J., & Neubig, G. (2020). How Can We Know What Language Models Know? *Transactions of the Association for Computational Linguistics*, 8, 423–438.

https://doi.org/10.1162/tac1_a_00324

Kotelnikova, A. V., Vychezhnanin, S. V., & Kotelnikov, E. V. (2023). Cross-Domain Sentiment Analysis Based on Small in-Domain Fine-Tuning. *IEEE Access : Practical Innovations, Open Solutions*, 11, 41061–41074. <https://doi.org/10.1109/access.2023.3269720>

Kumar, A., & Albuquerque, V. H. C. (2021). Sentiment Analysis Using XLM-R Transformer and Zero-shot Transfer Learning on Resource-poor Indian Language. *ACM Transactions on*

- Asian and Low-Resource Language Information Processing*, 20(5), 1–13.
<https://doi.org/10.1145/3461764>
- Laurer, M., van Atteveldt, W., Casas, A., & Welbers, K. (2023). Less Annotating, More Classifying: Addressing the Data Scarcity Issue of Supervised Machine Learning with Deep Transfer Learning and BERT-NLI. *Political Analysis*, 32(1), 84–100.
<https://doi.org/10.1017/pan.2023.20>
- Lin, X. V., Mihaylov, T., Artetxe, M., Wang, T., Chen, S., Simig, D., Ott, M., Goyal, N., Bhosale, S., Du, J., Pasunuru, R., Shleifer, S., Koura, P. S., Chaudhary, V., O'Horo, B., Wang, J., Zettlemoyer, L., Kozareva, Z., Diab, M., ... Li, X. (2022). Few-shot Learning with Multilingual Generative Language Models. *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*.
<https://doi.org/10.18653/v1/2022.emnlp-main.616>
- Michailidis, P. D. (2024). A Comparative Study of Sentiment Classification Models for Greek Reviews. *Big Data and Cognitive Computing*, 8(9), 107.
<https://doi.org/10.3390/bdcc8090107>
- Mikhail, K., O., K., Liliana, C. H., Grigori, S., & A., G. (2024). Multilingual Approaches to Sentiment Analysis of Texts in Linguistically Diverse Languages: A Case Study of Finnish, Hungarian, and Bulgarian. *IWCLUL*.
- Muennighoff, N., Wang, T., Sutawika, L., Roberts, A., Biderman, S., Le Scao, T., Bari, M. S., Shen, S., Yong, Z. X., Schoelkopf, H., Tang, X., Radev, D., Aji, A. F., Almubarak, K., Albanie, S., Alyafeai, Z., Webson, A., Raff, E., & Raffel, C. (2023). Crosslingual Generalization through Multitask Finetuning. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.
<https://doi.org/10.18653/v1/2023.acl-long.891>
- Nazir, M. K., Faisal, C. N., Habib, M. A., & Ahmad, H. (2025). Leveraging Multilingual Transformer for Multiclass Sentiment Analysis in Code-Mixed Data of Low-Resource Languages. *IEEE Access : Practical Innovations, Open Solutions*, 13, 7538–7554.

- <https://doi.org/10.1109/access.2025.3527710>
- Nie, E., Yuan, S., Ma, B., Schmid, H., Färber, M., Kreuter, F., & Schütze, H. (2024). Decomposed Prompting: Unveiling Multilingual Linguistic Structure Knowledge in English-Centric Large Language Models. *arXiv.org*. <https://doi.org/10.48550/ARXIV.2402.18397>
- Pan, R., García-Díaz, J. A., Garcia-Sanchez, F., & Valencia-García, R. (2023). Evaluation of transformer models for financial targeted sentiment analysis in Spanish. *PeerJ Computer Science*, 9, e1377. <https://doi.org/10.7717/peerj-cs.1377>
- Perron, B. E., Luan, H., Qi, Z., Victor, B. G., & Goyal, K. (2024). Demystifying Application Programming Interfaces (APIs): Unlocking the Power of Large Language Models and Other Web-based AI Services in Social Work Research. *arXiv.org*. <https://doi.org/10.48550/ARXIV.2410.20211>
- Plagwitz, L., Neuhaus, P., Yildirim, K., Losch, N., Varghese, J., & Büscher, A. (2024). Zero-Shot LLMs for Named Entity Recognition: Targeting Cardiac Function Indicators in German Clinical Texts. In *Studies in Health Technology and Informatics*. IOS Press. <https://doi.org/10.3233/shti240861>
- Příbáň, P., Šmíd, J., Steinberger, J., & Mištera, A. (2024). A comparative study of cross-lingual sentiment analysis. *Expert Systems with Applications*, 247, 123247. <https://doi.org/10.1016/j.eswa.2024.123247>
- Rafeef A. Hameed, Wael J. Abed, & Ahmed T. Sadiq. (2023). Evaluation of Hotel Performance with Sentiment Analysis by Deep Learning Techniques. *International Journal of Interactive Mobile Technologies (iJIM)*, 17(09), 70–87. <https://doi.org/10.3991/ijim.v17i09.38755>
- Rusnachenko, N., Golubev, A., & Loukachevitch, N. (2024). Large Language Models in Targeted Sentiment Analysis. *arXiv.org*. <https://doi.org/10.48550/ARXIV.2404.12342>
- Salahudeen, S. A., Lawan, F. I., Wali, A. M., Imam, A. A., Shuaibu, A. R., Yusuf, A., Rabi, N. B., Bello, M., Adamu, S. U., Aliyu, S. M., Gadanya, M. S., Muaz, S. A., Ahmad, M. S., Abdullahi, A., & Jamoh, A. Y. (2023). HausaNLP at SemEval-2023 Task 12: Leveraging

- African Low Resource TweetData for Sentiment Analysis. *International Workshop on Semantic Evaluation*. <https://doi.org/10.48550/ARXIV.2304.13634>
- Strubell, E., Ganesh, A., & McCallum, A. (2019). Energy and Policy Considerations for Deep Learning in NLP. *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. <https://doi.org/10.18653/v1/p19-1355>
- Tela, A., Woubie, A., & Hautamäki, V. (2024). Transferring monolingual model to low-resource language: The case of Tigrinya. *Applied Computing and Intelligence*, 4(2), 184–194. <https://doi.org/10.3934/aci.2024011>
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., Rodriguez, A., Joulin, A., Grave, E., & Lample, G. (2023). *LLaMA: Open and Efficient Foundation Language Models* (arXiv:2302.13971). arXiv. <https://arxiv.org/abs/2302.13971>
- Venerito, V., & Iannone, F. (2024). Large language model-driven sentiment analysis for facilitating fibromyalgia diagnosis. *RMD Open*, 10(2), e004367. <https://doi.org/10.1136/rmdopen-2024-004367>
- Vileikytė, B., Lukoševičius, M., & Stankevičius, L. (2024). Sentiment Analysis of Lithuanian Online Reviews Using Large Language Models. *International Conference on Information Technology*. <https://doi.org/10.48550/ARXIV.2407.19914>
- Young, L., & Soroka, S. (2012). Affective News: The Automated Coding of Sentiment in Political Texts. *Political Communication*, 29(2), 205–231. <https://doi.org/10.1080/10584609.2012.671234>

Appendix A

Appendix B

Figure B1
Sentiment Analysis Models Mean Absolute Error




Figure 4 shows the Mean Absolute Error (MAE) results across models and conditions, providing complementary insights into prediction accuracy beyond correlation analysis.

Figure B2
Sentiment Analysis Comparison Over Time




Figure 5 compares corpus-wide temporal sentiment trends captured by the best-performing LLM and by the Lexicoder Sentiment Dictionary, illustrating that the two methods track closely despite the LLM's higher correlation with sentence-level manual annotations.

Table B1*F1 Scores for 3-Category Classification*

Model	negative	neutral	positive	weighted_f1
deepseekchat_fr_fr	0.768	0.35	0.828	0.78
deepseekchat_en_fr	0.772	0.364	0.822	0.778
claude35_en_fr	0.748	0.408	0.804	0.762
mistral_en_en	0.659	0.414	0.789	0.723
claude35_fr_fr	0.742	0.353	0.749	0.723
deepseekchat_en_en	0.714	0.444	0.735	0.711
mistral_en_fr	0.612	0.273	0.805	0.71
gemma29b_fr_fr	0.685	0.262	0.756	0.703
gemma29b_en_fr	0.707	0.301	0.738	0.702
llama323b_fr_fr	0.637	0.105	0.756	0.678
llama323b_en_en	0.578	0.235	0.758	0.667
mistral_fr_fr	0.615	0.111	0.742	0.662
claude35_en_en	0.707	0.222	0.642	0.639
gemma29b_en_en	0.667	0.247	0.655	0.635
gemini20_en_fr	0.733	0.226	0.62	0.634
llama3370b_fr_fr	0.691	0.211	0.613	0.615
gemini20_fr_fr	0.707	0.214	0.601	0.613
llama3370b_en_en	0.637	0.195	0.607	0.593
llama323b_en_fr	0.308	0.308	0.766	0.587
gpt4o_en_fr	0.713	0.211	0.553	0.586
llama3370b_en_fr	0.657	0.202	0.584	0.586
lsd_en	0.641	0.192	0.594	0.585
gpt4o_fr_fr	0.709	0.203	0.532	0.571
gpt4o_en_en	0.703	0.19	0.507	0.553
lsd_fr	0.444	0.245	0.594	0.524
llama321b_en_en	0.459	0	0.561	0.494
deepseekr1_en_en	0	0.111	0.755	0.467
qwq32b_en_en	0	0	0.76	0.464
qwq32b_fr_fr	0	0	0.758	0.462
qwq32b_en_fr	0	0	0.757	0.462
deepseekr1_en_fr	0	0.261	0.73	0.461
deepseekr1_fr_fr	0	0.111	0.726	0.45
gemini20_en_en	0.508	0.147	0.43	0.439
llama321b_en_fr	0.358	0	0.478	0.41
llama321b_fr_fr	0.352	0.067	0.381	0.352

Table B2*F1 Scores for 7-Category Classification*

Model	very_neg	negative	somewhat_neg	neutral	somewhat_pos	positive	very_pos	weighted
claude35-en-fr	0.2	0.316	0.449	0.408	0.476	0.515	0.286	0.451
deepseekchat-fr-fr	0.222	0.431	0.37	0.35	0.541	0.364	0.062	0.448
deepseekchat-en-fr	0.182	0.467	0.421	0.364	0.464	0.36	0.067	0.428
claude35-fr-fr	0.267	0.476	0.46	0.353	0.348	0.6	0	0.426
deepseekchat-en-en	0.182	0.406	0.308	0.444	0.413	0.488	0.154	0.402
mistral-en-en	0.364	0.415	0.24	0.414	0.405	0.5	0	0.386
mistral-en-fr	0.222	0.423	0.302	0.273	0.349	0.452	0.154	0.358
claude35-en-en	0.286	0.5	0.465	0.222	0.263	0.417	0.222	0.355
llama323b-en-en	0.235	0.423	0.203	0.235	0.379	0.453	0	0.349
llama323b-fr-fr	0	0.467	0.179	0.105	0.4	0.43	0	0.344
llama321b-en-en	0	0.16	0.368	0	0.455	0.194	0	0.327
gemma29b-en-fr	0.286	0.519	0.107	0.301	0.31	0.429	0.077	0.31
gemini20-en-fr	0.444	0.579	0.246	0.226	0.157	0.639	0	0.309
gemma29b-fr-fr	0.267	0.526	0.069	0.262	0.306	0.426	0.083	0.298
mistral-fr-fr	0.4	0.424	0.28	0.111	0.262	0.382	0.167	0.296
gemini20-fr-fr	0.333	0.514	0.215	0.214	0.15	0.565	0.25	0.28
lsd-en	0	0	0.384	0.192	0.425	0	0	0.278
gemma29b-en-en	0.25	0.42	0.169	0.247	0.232	0.442	0.1	0.276
lsd-fr	0	0	0.248	0.245	0.447	0.049	0	0.271
gpt4o-en-fr	0.333	0.588	0.077	0.211	0.122	0.636	0	0.257
gemini20-en-en	0	0.453	0.215	0.147	0.182	0.431	0	0.257
gpt4o-fr-fr	0.444	0.595	0.069	0.203	0.105	0.644	0	0.251
gpt4o-en-en	0.444	0.632	0.2	0.19	0.096	0.457	0	0.247
llama323b-en-fr	0	0.108	0.068	0.308	0.3	0.375	0	0.239
llama3370b-en-fr	0.174	0.423	0.237	0.202	0.073	0.505	0.2	0.227
llama3370b-fr-fr	0.211	0.444	0.111	0.211	0.108	0.533	0.333	0.225
llama321b-en-fr	0.25	0.209	0.206	0	0.301	0.131	0	0.223
llama321b-fr-fr	0	0.103	0.288	0.067	0.277	0.109	0	0.216
llama3370b-en-en	0.211	0.456	0.169	0.195	0.036	0.416	0.182	0.186
deepseekr1-en-fr	0	0	0	0.261	0.105	0.359	0	0.121
deepseekr1-fr-fr	0	0	0	0.111	0.157	0.176	0.013	0.105
qwq32b-en-fr	0	0	0	0	0.019	0.114	0.01	0.027
deepseekr1-en-en	0	0	0	0.111	0	0.122	0.009	0.027